

# Introduction to Concurrent Programming

David Marchant

Based on slides by Troels Henriksen

2023-10-09

Inspired by slides by Randal E. Bryant and David R. O'Hallaron.

Why concurrency is hard and sometimes useful

The Unix model of concurrency

Memory model for threads

Mutexes

Scalability

# What is concurrent programming?

From a programming perspective, *concurrency* means multiple *logical control flows* executing simultaneously.

## Logical control flow

A stream of execution where the choice of what to do next is made by the code itself (i.e. this is pretty much all code you've written so far).

- **Concurrency is almost always present.**
  - ▶ E.g. the code rendering this slide runs *concurrently* with various system-level maintenance tasks, or perhaps a web browser.
  - ▶ These concurrent processes are isolated and do very different things.
  - ▶ Gets interesting when multiple control flows *interact*, e.g. by modifying shared data.

# We need to use words deliberately

Before we even get into any programming or background, we need to establish two **key** words.

## Parallel

Two or more logical flows of events happening or existing at literally the same time.

## Concurrent


Two or more logical flows of events that may happen in any order, and even be interleaved.

# We have already used these concepts

- We have already seen this with processes (see 2-3 lectures ago).
- We can have multiple processes executing concurrently, where part of one is progressing, then part of another.
- We can have two physical processors operating in parallel, each at literally the same time.
- We can execute concurrently on a single processor, not in parallel. But 2 processors means we can execute 2 things in parallel.

# The core idea

- This bit ain't complicated
- We have  $n$  many resources, not just one. We should be able to do  $n$  many things at once
- For instance a 4 core laptop should be doing 4 things in *parallel* (many more *concurrently*)



As we will frequently see, this idealised result is seldom possible.

# High level example of shared state

**A**

```
1 x = 1;  
2 y = 2;  
3 x = y + x;
```

**B**

```
1 x = 2;  
2 y = 1;  
3 x = y - x;
```

- If **A** runs first, then **B**:  $x=-1, y=1$ .
- If **B** runs first, then **A**:  $x=3, y=2$ .
- But any interleaving is also possible:
  - ▶ **A1, B1, B2, B3, A2, A3**:  $x=4, y=2$
- Ordering is preserved *within* each control flow, but unpredictable across control flows due to scheduling.

We will return *frequently* to the issue of synchronisation and nondeterministic execution.

# Motivation for concurrency: modularity

Sometimes multiple control flows is a natural way to express computation.

- Examples**
- **Video games:** distinct control flows for
    - ▶ Rendering graphics
    - ▶ Computing physics
    - ▶ AI for actors
    - ▶ Multiplayer communication
  - **Browsers**
    - ▶ A control flow per tab.
    - ▶ Maybe a control flow per resource (images etc) when downloading page.
  - **Network servers**
    - ▶ Control flow per user request.
- Note**
- This is useful even on small or old single-core processors.



# Motivation for concurrency: performance

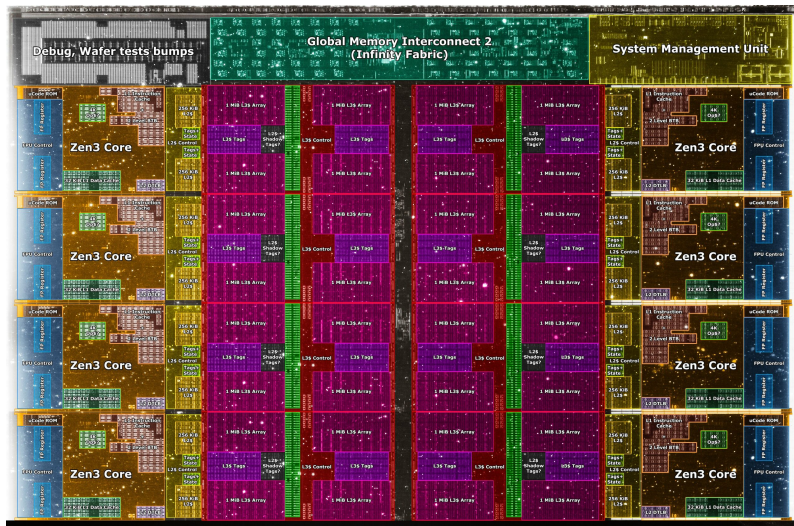
A single control flow occupies only a single *CPU core*.

## CPU core

A piece of hardware that executes instructions—contains registers and one or more levels of cache.

- When we previously talked about “the CPU” we really meant “a CPU core”.
- A modern CPU has several of these.
- **Each core executes a single logical control flow.**
- If we want to *fully utilize* the CPU (meaning: go fast), we need *a control flow per core*.

# AMD Ryzen 5000 (Zen 3 architecture)



<https://wccfttech.com/amd-ryzen-5000-zen-3-vermeer-undressed-high-res-die-shots-close-ups-pictured-detailed/>

# Problem: concurrent programming is hard!

- The human mind is sequential.
- Thinking about all possible orderings of events in a concurrent system is at least error prone and usually impossible.

# Problem: concurrent programming is hard!

- The human mind is sequential.
- Thinking about all possible orderings of events in a concurrent system is at least error prone and usually impossible.



The human brain is almost unchanged since this was the most complex problem it had to solve.

# Classic problems of concurrent programming

**Races:** Outcome depends on arbitrary scheduling decisions elsewhere.

- **Example:** who gets the last seat on the airplane?

# Classic problems of concurrent programming

**Races:** Outcome depends on arbitrary scheduling decisions elsewhere.

- **Example:** who gets the last seat on the airplane?

**Deadlock:** Resource allocation prevents forward progress.

- **Example:** traffic gridlock. (And programs generally cannot reverse!)

# Classic problems of concurrent programming

**Races:** Outcome depends on arbitrary scheduling decisions elsewhere.

- **Example:** who gets the last seat on the airplane?

**Deadlock:** Resource allocation prevents forward progress.

- **Example:** traffic gridlock. (And programs generally cannot reverse!)

**Starvation:** External events or scheduling prevents forward progress.

- **Example:** someone always jumping ahead in line.
- Also known as *livelock* or *fairness*.

# Classic problems of concurrent programming

**Races:** Outcome depends on arbitrary scheduling decisions elsewhere.

- **Example:** who gets the last seat on the airplane?

**Deadlock:** Resource allocation prevents forward progress.

- **Example:** traffic gridlock. (And programs generally cannot reverse!)

**Starvation:** External events or scheduling prevents forward progress.

- **Example:** someone always jumping ahead in line.
- Also known as *livelock* or *fairness*.

**This is a field that has more terms for how things can go *wrong* than go *right*.**

- By far the most difficult form of programming I know of.
- Our own brain is poorly suited for this kind of thinking.
  - ▶ We shall see C is not much better.
- Many aspects are outside the scope of CompSys.
  - ▶ But not all!



Why concurrency is hard and sometimes useful

The Unix model of concurrency

Memory model for threads

Mutexes

Scalability

# We have already seen multiple control flows

- *Processes* are an example of concurrent execution.
  - ▶ Generally *just work* because they are isolated from each other.
  - ▶ Interleaved execution often not noticeable.
  - ▶ Process isolation makes close collaboration inefficient and awkward.
- Instead we use *threads*.

# Our traditional model of a process

**A process consists of three parts:**

1. A *virtual memory space*.
  - ▶ Contains stack, code, heap, data, etc.
2. A *kernel context*.
  - ▶ PID, open files, signal mask, parent PID, list of children, etc.
3. An *execution context*.
  - ▶ Registers (including special ones like the program counter).

# A multi-threaded process model

**A process still consists of three parts:**

1. A *virtual memory space*.
  - ▶ Contains stack, code, heap, data, etc.
2. A *kernel context*.
  - ▶ Process ID (PID), open files, signal mask, parent PID, list of children, etc.
3. One or more *threads*, each containing.
  - 3.1 An *execution context*
    - ▶ Registers (including special ones like the program counter).
  - 3.2 A *kernel thread context*.
    - ▶ Thread ID (TID), and a few other things that do not matter.

**Processes consist of one or more threads!**

# Threads and sharing

## Threads in the same process

- Each thread has its own logical control flow.
  - Each thread has its own stack.
  - Threads share open files.
  - Threads share the same virtual memory space.
  - They are *peers*, there is no “main thread”.
- 
- **Implication:** threads can interact by modifying memory.
    - ▶ Even unintentionally.

# Threads contra processes

## Similarities

- Each has its own logical control flow.
- Each can run concurrently with others (possibly also parallel).
- Each is context switched.

## Differences

- Threads share code and data.
  - ▶ Processes typically do not.
- Threads are cheaper to create and maintain than processes.
  - ▶ Take with a grain of salt; both are plenty fast on Linux.
  - ▶ But switching between threads within a process does not require switching to a new virtual address space.
- Each process has one or more threads.
- Each thread belongs to exactly one process.

# POSIX threads—standard thread interface on Unix

- Creating and reaping threads:
  - ▶ `pthread_create()`
  - ▶ `pthread_join()`
- Determining your thread ID:
  - ▶ `pthread_self()`
- Terminating threads:
  - ▶ `pthread_cancel()` (using this is usually a mistake)
  - ▶ `pthread_exit()` – terminates calling thread.
  - ▶ `exit()` – terminates *all* threads.
    - ▶ Implicit when `main()` returns.
- Synchronisation:
  - ▶ `pthread_mutex_init()`
  - ▶ `pthread_mutex_lock()`
  - ▶ `pthread_mutex_unlock()`

We will add a few more functions next lecture, but this is plenty to get in trouble.

# Hello World in POSIX threads



# Hello World in POSIX threads

```
#include <pthread.h>
#include <assert.h>
#include <stdio.h>

void* thread(void *arg) {
    int* p = arg;
    printf("Hello world! %d\n", *p);
    return NULL;
}

int main(void) {
    int x = 42;
    pthread_t tid;
    assert(pthread_create(&tid, NULL, thread, &x) == 0);
    assert(pthread_join(tid, NULL) == 0);
}
```

Why concurrency is hard and sometimes useful

The Unix model of concurrency

Memory model for threads

Mutexes

Scalability

## Example program to illustrate sharing

```
char **ptr;
int cnt = 0;
int main() {
    pthread_t tid;
    char *msgs[2] = {
        "Hello_from_foo",
        "Hello_from_bar"
    };
    ptr = msgs;
    for (int i = 0; i < 2; i++)
        pthread_create(&tid,
                       NULL,
                       thread,
                       (void *)i);
    pthread_exit(NULL);
}
```

```
void* thread(void *vargp) {
    int j = (int)vargp;

    printf("%d:_%s_(cnt=%d)\n",
           j, ptr[j], ++cnt);
    return NULL;
}
```

- **Global variables**—*one instance.*
- **Local variables**—*one instance per function call.*
- Variables are shared if multiple threads reference the same instance.
- **Which variables are shared here?**

## Example program to illustrate sharing

```
char **ptr;
int cnt = 0;
int main() {
    pthread_t tid;
    char *msgs[2] = {
        "Hello_from_foo",
        "Hello_from_bar"
    };
    ptr = msgs;
    for (int i = 0; i < 2; i++)
        pthread_create(&tid,
                       NULL,
                       thread,
                       (void *)i);
    pthread_exit(NULL);
}
```

```
void* thread(void *vargp) {
    int j = (int)vargp;

    printf("%d:_%s_(cnt=%d)\n",
           j, ptr[j], ++cnt);
    return NULL;
}
```

- **Global variables**—*one instance*.
- **Local variables**—*one instance per function call*.
- Variables are shared if multiple threads reference the same instance.
- **Which variables are shared here?**
  - ▶ ptr, cnt, msgs.

# Synchronising threads

```
int n = atoi(argv[1]);
pthread_t tid1, tid2;
pthread_create(&tid1,
               NULL,
               thread,
               &n);
pthread_create(&tid2,
               NULL,
               thread,
               &n);
pthread_join(tid1, NULL);
pthread_join(tid2, NULL);
if (cnt != 2 * n)
    printf("Bad:_%d\n", cnt);
else
    printf("Good:_%d\n", cnt);
}
```

```
int cnt = 0;

void *thread(void *vargp) {
    int n = *((int*)vargp);

    for (int i = 0; i < n; i++)
        cnt++;

    return NULL;
}
```

- Updates of shared variable `cnt` result in nondeterminism.
- But why?

# Assembly code for loop

```
for (int i = 0; i < n; i++)  
    cnt++;
```

```
    H {  
        li t0, 0           # i = 0  
        la t1, cnt         # address of cnt in t1  
        beq t0, a0, done   # skip if nothing to do  
loop:  
    L {  
        lw t2, 0(t1)       # load cnt from memory  
    U {  
        addi t2, t2, 1      # increment cnt  
    S {  
        sw t2, 0(t1)       # store cnt in memory  
    T {  
        addi t0, t0, 1      # i++  
        beq t0, a0, done   # done?  
        j loop             # another iteration  
done:
```

# Concurrent execution, when we are lucky

- Any sequentially consistent interleaving is possible, and some give an unexpected result.

$i$	instruction	$t2_0$	$t2_1$	cnt
0	$H_0$	?	?	0
0	$L_0$	0	?	0
0	$U_0$	1	?	0
0	$S_0$	1	?	1
1	$H_1$	1	1	1
1	$L_1$	1	1	1
1	$U_1$	1	2	1
1	$S_1$	1	2	2
1	$T_1$	1	2	2
0	$T_0$	1	2	2



Thread 0 critical section  
Thread 1 critical section

Correct result!

# Concurrent execution, when we are not so lucky

- Any sequentially consistent interleaving is possible, and some give an unexpected result.

$i$	instruction	$t2_0$	$t2_1$	cnt
0	$H_0$	?	?	0
0	$L_0$	0	?	0
0	$U_0$	1	?	0
1	$H_1$	1	0	0
1	$L_1$	1	0	0
0	$S_0$	1	0	1
0	$T_0$	1	1	1
1	$U_1$	1	1	1
1	$S_1$	1	1	1
1	$T_1$	1	1	1



Thread 0 critical section  
Thread 1 critical section

Wrong result!



# Critical sections and atomicity

## Definition (Critical section)

A section of code that only a single thread must be executing at a time.

- The general principle is **mutual exclusion**.
- How do we ensure that critical sections are executed atomically?
- **Mutexes!**

Why concurrency is hard and sometimes useful

The Unix model of concurrency

Memory model for threads

**Mutexes**

Scalability

# Mutex API

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

- `pthread_mutex_t` is the pthreads *mutex type*.
- Mutexes have two states: *locked* and *unlocked*
- New mutexes start out *unlocked*.

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Locking an a currently *unlocked* mutex *locks it* and returns.
  - ▶ The thread now *holds* the mutex.
- Trying to lock a currently *locked* mutex **blocks the calling thread**.
  - ▶ The thread now *waits* for the mutex.
  - ▶ When the mutex is unlocked (by some other thread), one *waiting* thread is allowed to lock the mutex.
- **Strong property:** when `pthread_mutex_lock()` returns (assuming no error), *the calling thread holds the mutex*.

# Basic idea: Protecting critical sections with mutexes

- We associate each critical section with a mutex.
- **When entering the critical section:** lock the mutex.
- **When leaving the critical section:** unlock the mutex.

```
// shared mutex instance  
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
pthread_mutex_lock(&mutex);  
... critical section ...  
pthread_mutex_unlock(&mutex);
```

# Basic idea: Protecting critical sections with mutexes

- We associate each critical section with a mutex.
- **When entering the critical section:** lock the mutex.
- **When leaving the critical section:** unlock the mutex.

```
// shared mutex instance  
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
pthread_mutex_lock(&mutex);  
... critical section ...  
pthread_mutex_unlock(&mutex);
```

- C doesn't have a language-level notion of “critical section”, and typically it isn't *code* that we need to protect.
- In practice we associate mutexes with *shared variables*.
  - ▶ Sometimes a single mutex protects *multiple* variables.

# Protecting the critical section in the counting program

```
pthread_mutex_t cnt_mutex = PTHREAD_MUTEX_INITIALIZER;  
int cnt = 0;
```

```
for (int i = 0; i < n; i++) {  
    pthread_mutex_lock(&cnt_mutex);  
    cnt++;  
    pthread_mutex_unlock(&cnt_mutex);  
}
```

# Gotta go fast

```
int local_cnt = 0;
for (int i = 0; i < n; i++) {
    local_cnt++;
}

pthread_mutex_lock(&cnt_mutex);
cnt += local_cnt;
pthread_mutex_unlock(&cnt_mutex);
```

# How mutexes are implemented

Real mutexes contain more features, but this captures the essence:

```
struct mutex {  
    int locked; // 0 if unlocked, 1 if locked.  
};
```

```
void lock(struct mutex *mutex) {  
    while (mutex->locked != 0) {  
        // try again  
    }  
    mutex->locked = 1;  
}
```

**Problem?**



# How mutexes are implemented

Real mutexes contain more features, but this captures the essence:

```
struct mutex {  
    int locked; // 0 if unlocked, 1 if locked.  
};
```

```
void lock(struct mutex *mutex) {  
    while (mutex->locked != 0) {  
        // try again  
    }  
    mutex->locked = 1;  
}
```

**Problem?**

Not atomic.

# The compare-and-swap (CAS) operation

```
int cas(int* p, int expected, int desired) {  
    // This should be atomic!  
    if (*p == expected) {  
        *p = desired;  
        return 0;  
    } else {  
        return 1;  
    }  
}  
  
void lock(struct mutex *mutex) {  
    while (1) {  
        // try to switch from 0 to 1  
        int switched = cas(&mutex->locked, 0, 1);  
        if (switched) { break; }  
    }  
}
```

OK, but we still have the same problem.

# Compare-and-swap

```
int cas(int* p, int expected, int desired) {  
    // This should be atomic!  
    if (*p == expected) {  
        *p = desired;  
        return 0;  
    } else {  
        return 1;  
    }  
}
```

- This function cannot be implemented in C.
- In fact, is only possible if the architecture provides it as a primitive (as x86 does), or some even more basic primitive (RISC-V *conditional loads/stores*).

# Compare-and-Swap in RISC-V

```
# a0 holds address of memory location
# a1 holds expected value
# a2 holds desired value
cas:  lr.w t0, (a0)      # load original value from a0 into t0
      bne t0, a1, fail  # value != expected, so fail
      sc.w t0, a2, (a0) # if (a0) untouched, write a2 to a0
      bnez t0, fail     # store-conditional failed, so CAS failed
      li  a0, 0         # set return to success
      jr  ra            # return
fail:  li  a0, 1         # set return to failure
      jr  ra            # return
```

- Real implementations are actually a little more complicated.
- Students who wish to derail their life can study the details at <https://five-embeddev.com/riscv-isa-manual/latest/a.html>

Why concurrency is hard and sometimes useful

The Unix model of concurrency

Memory model for threads

Mutexes

Scalability

# Speedup

## Definition (Speedup in latency)

If  $T_1, T_2$  are the runtimes of two programs  $P_1, P_2$ , then the *speedup in latency* of  $P_2$  over  $P_1$  is

$$\frac{T_1}{T_2}$$

## Definition (Speedup in throughput)

If  $Q_1, Q_2$  are the throughputs of two programs  $P_1, P_2$ , then the *speedup in throughput* of  $P_2$  over  $P_1$  is

$$\frac{Q_2}{Q_1}$$

# Scalability

## Definition (Strong scaling)

How the runtime varies with the number of processors for a fixed problem size.

## Definition (Weak scaling)

How the runtime varies with the number of processors for a fixed problem size *relative to the number of processors*.

## Definition (Amdahl's Law)

If  $p$  is the proportion of execution time that benefits from parallelisation, then  $S(N)$  is maximum theoretical speedup achievable by execution on  $N$  threads, and is given by

$$S(N) = \frac{1}{(1 - p) + \frac{p}{N}}$$


Note that

$$S(N) \leq \frac{1}{1 - p}$$

- Potential speedup by optimising part of system is bounded by proportion of part in overall runtime.
- **We should optimise the parts that take a long while to run.**
- Predicts *strong scaling*.



# Amdahl's law is pessimistic



# Gustafson's Law


## Definition (Gustafson's Law)

If  $s = 1 - p$  is the proportion of execution time that must be sequential, then  $S(N)$  is maximum theoretical speedup achievable by execution on  $N$  threads, and is given by

$$S(N) = N + (1 - N) \times s$$

- Predicts *weak scaling*.

# Gustafson's law is optimistic



# Summary

- Concurrency has two goals: modularity and parallel execution (performance).
  - ▶ The latter is increasingly important.
- Concurrency causes nondeterministic execution.
  - ▶ *Hard* to reason about.
  - ▶ The machine is certain to betray you.
- Use mutexes to ensure *mutual exclusion* to variables or critical sections.
  - ▶ Ultimately based on tiny operations that the hardware guarantees are atomic.